

7. Feature Store

7.1 Feature Store Approach

A **custom feature store** is implemented to manage, document, and retrieve machine learning features in a consistent and versioned manner.

The feature store is designed to:

- Centralize access to engineered features
- Ensure consistency between training and inference
- Support feature versioning and metadata management

PostgreSQL is used as the **offline feature storage layer**, while a lightweight Python-based registry manages feature definitions and versions.

7.2 Feature Store Architecture

The feature store is built on top of the feature engineering pipeline and consists of the following components:



This architecture mirrors the design of production-grade feature stores while remaining simple and easy to maintain.

7.3 Feature Storage Layer

Engineered features are stored in PostgreSQL tables created during Step 6:

- `user_features`
- `item_features`
- `user_item_features`

These tables act as the **offline feature store**, enabling efficient querying and reuse across different stages of the ML pipeline.

7.4 Feature Metadata Registry

A centralized **feature registry** is used to document feature metadata and manage feature versions.

Registry Location

`src/feature_store/registry.py`

The registry captures:

- Feature group (user, item, user-item)
- Feature names
- Source tables
- Feature descriptions
- Feature version identifiers

This ensures feature discoverability and traceability.

7.5 Feature Versioning Strategy

Feature versions are managed using explicit version identifiers (e.g., `v1`).

Each version:

- Represents a stable set of feature definitions
- Can be independently retrieved for training or inference
- Enables backward compatibility when features evolve

Versioning guarantees reproducibility of model training and evaluation.

7.6 Feature Retrieval Mechanism

Feature retrieval is implemented using reusable Python functions that query PostgreSQL based on the requested feature version.

Retrieval functions support:

- User-level feature lookup
- Item-level feature lookup
- User-item interaction feature lookup

The same retrieval logic is used during **model training and inference**, ensuring feature consistency across the ML lifecycle.

7.7 Sample Feature Retrieval Demonstration

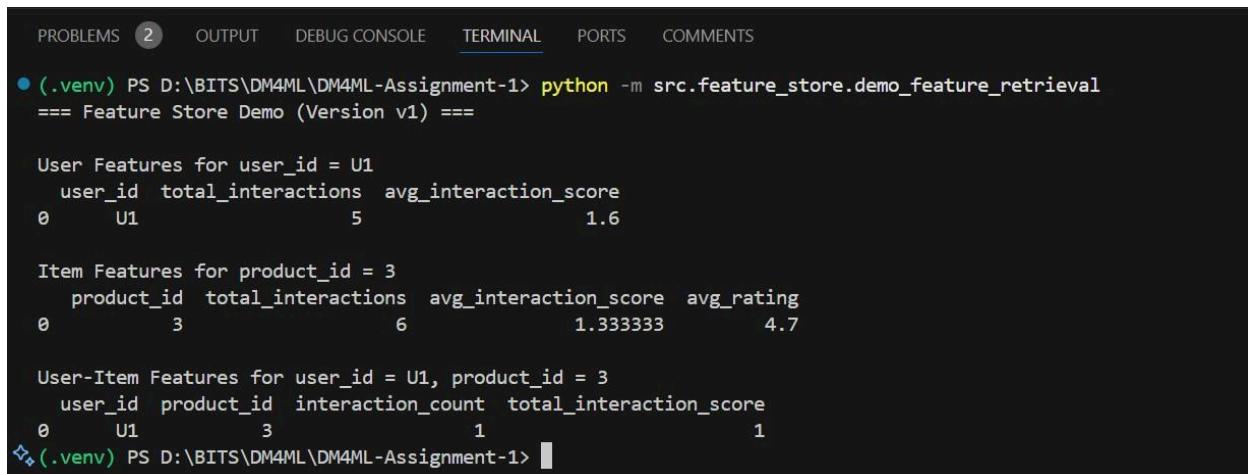
A standalone script demonstrates feature retrieval from the feature store.

Script Location

`src/feature_store/demo_feature_retrieval.py`

The script:

- Retrieves user features for a given user ID
- Retrieves item features for a given product ID
- Retrieves user-item interaction features
- Uses versioned feature access (**v1**)



A screenshot of a terminal window titled "TERMINAL". The window shows the command "python -m src.feature_store.demo_feature_retrieval" being run in a virtual environment named ".venv". The output displays three types of feature sets: User Features for user_id = U1, Item Features for product_id = 3, and User-Item Features for user_id = U1, product_id = 3. Each output is a table with columns like user_id, total_interactions, avg_interaction_score, etc.

```
● (.venv) PS D:\BITS\DM4ML\DM4ML-Assignment-1> python -m src.feature_store.demo_feature_retrieval
== Feature Store Demo (Version v1) ==

User Features for user_id = U1
  user_id  total_interactions  avg_interaction_score
  0        U1                  5                      1.6

Item Features for product_id = 3
  product_id  total_interactions  avg_interaction_score  avg_rating
  0            3                  6          1.333333           4.7

User-Item Features for user_id = U1, product_id = 3
  user_id  product_id  interaction_count  total_interaction_score
  0        U1          3                  1                      1

◆ (.venv) PS D:\BITS\DM4ML\DM4ML-Assignment-1>
```

This demonstrates how downstream ML components can consume features reliably.

7.8 Training and Inference Consistency

- Both training and inference pipelines access features through the same APIs
- Feature definitions and transformations remain synchronized
- Versioned retrieval prevents feature drift between environments

This ensures reliable and repeatable model behavior.

7.9 Summary

The feature store provides a structured, versioned, and documented mechanism for managing machine learning features. It bridges the gap between feature engineering and model consumption, enabling scalable and reproducible recommendation workflows.